

Threads

COMP 401 Fall 2018

Lecture 21

Threads

- As a generic term
 - Abstraction for program execution
 - Current point of execution.
 - Call stack.
 - Contents of memory.
 - The fundamental unit of processing that can be scheduled by an operating system.
- Multithreading
 - A program running two or more threads concurrently.
 - Separate points of execution.
 - Separate call stacks.
 - Shared memory.

Rise of Threads

- In the beginning was the command line...
 - Neal Stephenson
- First computers were batch uniprocessors
 - Ran one program at a time.
 - Non-interactive
 - No need for multithreading.

Time Division


- Computers getting faster/cheaper gave rise to interactive computing with GUIs.
 - More than one program at a time.
 - Processor time division
 - Operating system rapidly switches between separate processes.
 - Resource sharing.
 - Not quite multithreading
 - » Separate processes (i.e., programs)
 - » No memory sharing.
 - Illusion of private resource.

Threads For Time Division

- Threads extend OS mechanisms for process-level time division to within a program.
 - Program needs/wants to make progress on two or more tasks.
 - Similar to the idea of two separate processes sharing the processor to make “simultaneous” progress.
 - But in this case, two threads that are part of *same* program
 - Need to share memory and coordinate actions.
- Example: GUI
 - Want GUI to remain responsive when CPU heavy task is occurs.
 - lec21.v1


The Multicore Revolution

Microprocessor Transistor Counts 1971-2011 & Moore's Law



Moore's Law – The number of transistors on integrated circuit chips (1971-2016)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are strongly linked to Moore's law.



Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)


The data visualization is available at [OurWorldinData.org](https://ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under [CC-BY-SA](https://creativecommons.org/licenses/by-sa/4.0/) by the author Max Roser.


Transistors on integrated circuit chips (1971-2016)

Our World
in Data

Transistors on integrated circuits doubles approximately every two years.
Progress – such as processing speed or the price of electronic products – are



Core Count Driving Moore's Law



<https://madusudanan.com/blog/scala-tutorials-part-9-intro-to-functional-programming/>

Parallel Programming

- Multicore increases performance only if we can find ways to parallelize our task.
 - Sometimes very easy.
 - No data or logic dependencies.
 - Aggregating data operations.
 - Sometimes very hard.

Java Threads

- Two basic mechanisms for threads:
 - Extend Thread
 - Override run()
 - Call start()
 - Creates a new thread.
 - Executes run() in that thread.
 - lec21.v2
 - Note: only works once.
 - Implement Runnable
 - Provide reference to “runnable” object as parameter of Thread constructor.
 - Call start() on the thread object.
 - lec21.v3

Consistency

- Multi-threaded access to memory can be hazardous.
 - Operation ordering
 - May need all threads to see changes to memory in a consistent way.
 - Actual order may or may not matter as long as all threads see the same effective order.
 - Exclusive access
 - Operation may involve many steps which must be done with exclusive access to memory.
 - Other threads must be prevented from accessing memory while operation in progress.
- lec21.v4

Object Monitor

- Every object in Java associated with a “monitor”.
 - Provides ability to “lock” the object.
 - Only one thread can “own” the monitor at a time.
 - Strategy is to lock object before executing any unsafe code.
 - Two ways to acquire the monitor lock:
 - synchronized methods
 - synchronized statements

Synchronized Methods

- “synchronized” keyword
 - `public synchronized void a_method() {...`
- All synchronized methods obtain monitor lock of object before execution.
 - Automatically released when method returns.
 - This means that for all synchronized methods of an object, only one can be executing at any given time.
- Does not affect unsynchronized methods.
 - This means that all methods that could cause consistency problems must all be synchronized.

Synchronized Statements

- Syntax:

```
synchronized (object) {  
    ...  
}
```

- Lock obtained on monitor associated with object before statement block is executed.
- Allows synchronization to be wrapped around a small bit of code.
- Allows synchronization using object reference.

Synchronization Equivalence

- These two are equivalent:

```
public synchronized void a_method() {  
    ...  
}
```

```
public void a_method() {  
    synchronized(this) {  
        ...  
    }  
}
```


Consistency Example Revisited

- lec21.v5

Deadlock

- Synchronization can be tricky.
 - Consider:
 - Synchronized method in object A calling a synchronized method in object B.
 - Obtained lock on A
 - Waiting for lock on B
 - At same time in another thread, suppose synch. method of object B is calling synch. method in object A.
 - Obtained lock on B
 - Waiting for lock on A
 - Now neither thread can continue.
 - This is called deadlock.
- lec21.v6

Join

- Sometimes one thread needs to wait for another thread to be done before continuing.
 - A common form of thread coordination.
- `join()`
 - Method provided by Thread.
 - When called in a different thread, the calling thread is blocked until called thread is done.
 - Done = `run()` method has finished.
 - Blocked means `join()` does not return until called thread is done (thus calling thread is waiting).
- lec21.v7

wait() and notify()

- Mechanisms for coordination between running threads.
 - These are methods defined by Object
 - Must own the lock associated with the object in order to call any of these methods.
 - Effectively means that call to wait() / notify() must be within a synchronized method / statement.
 - wait()
 - Calling thread waits until notified.
 - notify()
 - Releases one waiting thread (as soon as lock is available)
 - notifyAll()
 - Releases all waiting threads
 - Each resumes in turn as lock becomes available.
- lec21.v8

sleep() and timing execution

- Thread.sleep(int ms)
 - Sleeps for ms milliseconds
 - Accuracy not guaranteed
- System.nanoTime()
 - Highest resolution timer available in Java
 - Return value is *long*
 - Difference in value between two different calls represents duration in nanoseconds.
 - Actual resolution is courser than a nanosecond.
- lec21.v9

Demonstrating Parallelism

- Generating 100 million random numbers
 - With varying number of threads from 1-16
 - How many cores does my machine have?